

What's new in PyTorch 1.0

Key changes researchers should be aware of

Contents

1. The Hybrid Frontend
2. Redesigned Distributed Package
3. Torch Hub
4. C++ Frontend (API unstable)
5. Bug fixes and deprecations
6. Performance improvements

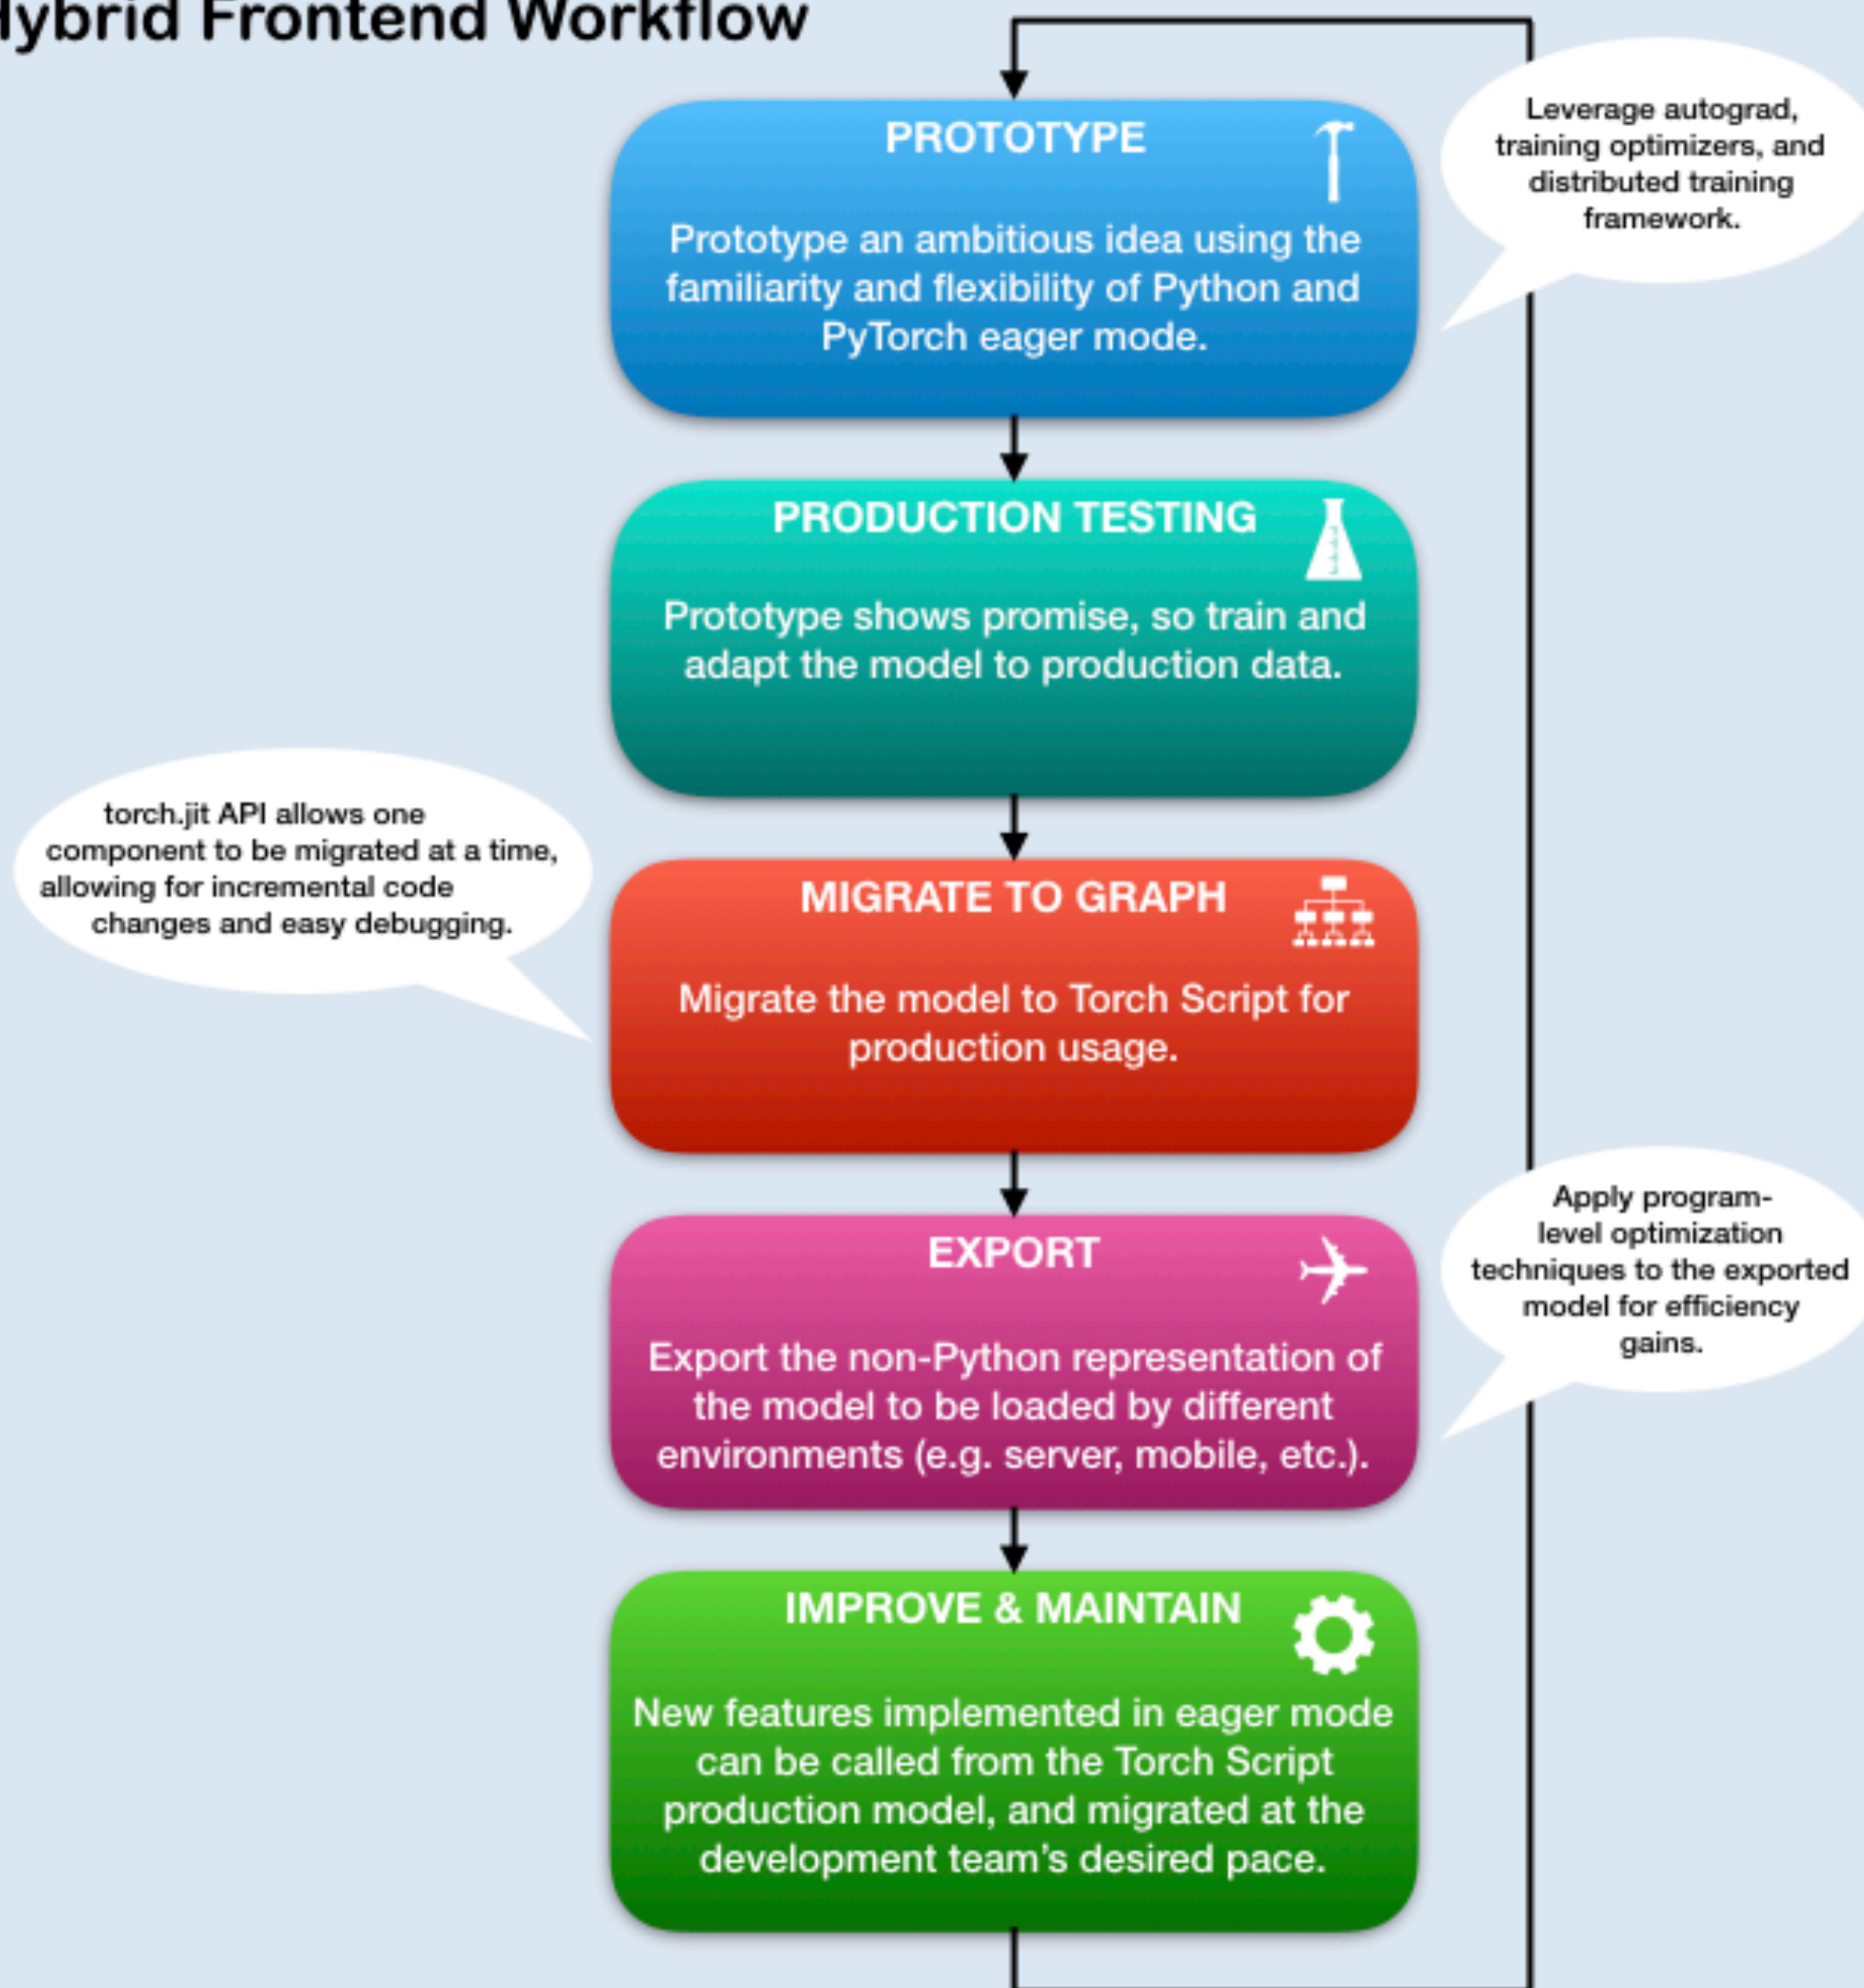
The Hybrid Frontend

- Graph-based model representation:
 - allows HW and SW optimizations
 - supports framework-agnostic model exports
 - suitable for production
- Imperative interface:
 - more Pythonic
 - easier debugging
 - suitable for research



torch.jit

Hybrid Frontend Workflow



The Hybrid Frontend

Tracing

- module or function + example inputs = graph-based function
- the model moves out of Python: no data-dependent control flow, no debugging

```
import torch
import torchvision

model = torchvision.models.resnet18()
example = torch.rand(1, 3, 224, 224)

traced_script_module = torch.jit.trace(model, example)
```

The Hybrid Frontend

Scripting

- Python → Torch Script: supports data-dependent control flow and debugging
- Torch Script does not support all Python features

```
class MyModule(torch.jit.ScriptModule):
    def __init__(self, N, M):
        super(MyModule, self).__init__()
        self.weight = torch.nn.Parameter(torch.rand(N, M))
    @torch.jit.script_method
    def forward(self, input):
        if input.sum() > 0:
            output = self.weight.mv(input)
```

The Hybrid Frontend

Related tutorials:

- [Loading a PyTorch Model in C++](#)
- [Deploying a seq2seq model with the Hybrid Frontend](#)
- [ONNX Live Tutorial](#)
- [Transferring a model from PyTorch to Caffe2 and mobile using ONNX](#)

Redesigned Distributed Package

- **torch.distributed**: base package with Gloo, NCCL, MPI backend support
- **torch.nn.parallel.DistributedDataParallel**: ready-to-go utility for data parallelism
- related tutorials:
 - [PyTorch 1.0 Distributed Trainer with Amazon Web Service](#)
 - [Writing Distributed Applications with PyTorch](#)

Torch Hub

Use GitHub as a hub for pretrained models!

- Define entrypoints in **hubconf.py**

```
dependencies = ['torch', 'math']

def resnet18(pretrained=False, *args, **kwargs):
    from torchvision.models.resnet import resnet18 as _resnet18

    model = _resnet18(*args, **kwargs)
    if pretrained:
        model.load_state_dict("pretrained.pth")

    return model
```

Torch Hub

- load entrypoints with **torch.hub.load**

```
hub_model = torch.hub.load(  
    'pytorch/vision:master',  
    'resnet18',  
    1234,  
    pretrained=True)
```

C++ Frontend (API unstable)

```
import torch

model = torch.nn.Linear(5, 1)
optimizer = torch.optim.SGD(model.parameters(), lr=0.1)
prediction = model.forward(torch.randn(3, 5))
loss = torch.nn.functional.mse_loss(prediction, torch.ones(3, 1))
loss.backward()
optimizer.step()
```

```
#include <torch/torch.h>

torch::nn::Linear model(5, 1);
torch::optim::SGD optimizer(model->parameters(), /*lr=*/0.1);
torch::Tensor prediction = model->forward(torch::randn({3, 5}));
auto loss = torch::mse_loss(prediction, torch::ones({3, 1}));
loss.backward();
optimizer.step();
```

Bug fixes and deprecations

- `torch.nn.functional.softmin` was using the incorrect formula in 0.4.1 (#10066)
- `DataLoader` fixed a couple of issues resulting in hangs (#11985, #12700)
- `torch.masked_fill_` now works properly on non-contiguous tensor inputs (#12594)
- `Tensor.__delitem__`: fixed a segmentation fault on (#12726)
- Removed support for C extensions. Please use cpp extensions (#12122)
- Deleted `torch.utils.trainer` (#12487)
- `torch.distributed.deprecated` and `torch.nn.parallel.deprecated.DistributedDataParallel` are deprecated but still available.

Performance improvements

- `torch.clamp` no longer does unnecessary copying (#10352)
- `torch.add`, `torch.sub`, `torch.mul`, `torch.div` are much faster for non-contiguous tensors on GPU (#8919)
- `torch.nn.RNN` and related Modules have been ported to C++ and are more performant (#10305, #10481)

Read the full release notes on
<https://github.com/pytorch/pytorch/releases/tag/v1.0.0>